

CS 61C Spring 2017 Discussion 10

MSI Cache Coherency

MSI is a simple concurrency protocol to prevent write-back caches from maintaining incorrect data. Its approach is to invalidate blocks other caches on write (instead of updating the value in the other caches).

State	Cache up to date?	Memory up to date?	Others have a copy?	Can respond to other's reads?	Can write without changing state?
Modified	Yes	No	No	Yes, Required	Yes
Shared	Yes	Maybe	Maybe	Yes, Optional	No
Invalid	No	Maybe	N/A	No	No

1. How many bits do we need to add to each cache block to implement MSI? **2**
2. Consider the following access pattern on a two-processor system with a memory size equivalent to two cache blocks. Each processor has a direct-mapped, write-back, write-allocate cache with only one cache block. Fill in the state of each cache, and in parentheses write the memory block contained in that cache.

Time	After Operation	P1 cache state	P2 cache state	Memory @ 0 up to date?	Memory @ 1 up to date?
0	P1: read block 1	Shared (1)	Invalid	YES	YES
1	P2: read block 1	Shared (1)	Shared (1)	YES	YES
2	P1: write block 1	Modified (1)	Invalid	YES	NO
3	P2: write block 1	Invalid	Modified (1)	YES	NO
4	P1: read block 0	Shared (0)	Modified (1)	YES	NO
5	P2: read block 0	Shared (0)	Shared (0)	YES	YES
6	P1: write block 0	Modified (0)	Invalid	NO	YES
7	P2: read block 0	Shared (0)	Shared (0)	NO	YES

Concurrency

1. Consider the following function:

```
void transferFunds(struct account *from, struct account *to, long cents) {  
    from->cents -= cents;  
    to->cents += cents;  
}
```

 - a. What are some data races that could occur if this function is called simultaneously from two (or more) threads on the same accounts? (Hint: if the problem isn't obvious, translate the function into MIPS first)
Each thread needs to read the "current" value, perform an add/sub, and store a value for from->cents and to->cents. Two threads could read the same "current" value and the later store essentially erases the other transaction at either line.
 - b. How could you fix or avoid these races? Can you do this without hardware support?
Wrap transferFunds in a critical section, or divide up the accounts array and for loop in a way that you can have separate threads work on different accounts

Thread Level Parallelism

```
#pragma omp parallelism
{
    /* code here */
}
```

*Each thread runs a copy of code within the block
*Thread scheduling is non-deterministic

```
#pragma omp parallel for
for (int i = 0; i < n; i++) {
    /* code here */
}
```

Same as:

```
#pragma omp parallel
{
    #pragma omp for
    for (int i = 0; i < n; i++) {...}
}
```

1. For the following snippets of code below, circle one of the following to indicate what issue, if any, the code will experience. Then provide a short justification. Assume the default number of threads is greater than 1. Assume no thread will complete before another thread starts executing. Assume arr is an int array with length n.

a)

// Set element i of arr to i

```
#pragma omp parallel
(int i = 0; i < n; i++)
    arr[i] = i;
```

Sometimes incorrect
Slower than serial

Always incorrect
Faster than serial

Slower than serial – there is no for directive, so every thread executes this loop in its entirety. n threads running n loops at the same time will actually execute in the same time as 1 thread running 1 loop. Despite the possibility of false sharing, the values should all be correct at the end of the loop. Furthermore, the existence of parallel overhead due to the extra number of threads could slow down the execution time.

b)

// Set arr to be an array of Fibonacci numbers.

```
arr[0] = 0;
arr[1] = 1;
#pragma omp parallel for
for (int i = 2; i < n; i++)
    arr[i] = arr[i-1] + arr[i-2];
```

Sometimes incorrect
Slower than serial

Always incorrect
Faster than serial

Always incorrect (if n>4) – Loop has data dependencies, so the calculation of all threads but the first one will depend on data from the previous thread. Because we said “assume no thread will complete before another thread starts executing,” then this code will always be wrong from reading incorrect values.

c)

// Set all elements in arr to 0;

```
int i;
#pragma omp parallel for
for (i = 0; i < n; i++)
    arr[i] = 0;
```

Sometimes incorrect
Slower than serial

Always incorrect
Faster than serial

Faster than serial – the for directive actually automatically makes loop variables (such as the index) private, so this will work properly. The for directive splits up the iterations of the loop into continuous chunks for each thread, so no data dependencies or false sharing.

2. Consider the following code:

```
// Decrements element i of arr. n is a multiple of omp_get_num_threads()
```

```
#pragma omp parallel {  
    int threadCount = omp_get_num_threads();  
    int myThread = omp_get_thread_num();  
    for (int i = 0; i < n; i++) {  
        if (i % threadCount == myThread)  
            arr[i] *= arr[i];  
    }  
}
```

What potential issue can arise from this code?

False sharing arises because different threads can modify elements located in the same memory block simultaneously. This is a problem because some threads may have incorrect values in their cache block when they modify the value `arr[i]`, invalidating the cache block. A fix to this will be discussed in lab.

Data Level Parallelism

<code>__m128i _mm_load1_si128()</code>	returns 128-bit one vector
<code>__m128i _mm_loadu_si128(__m128i *p)</code>	returns 128-bit vector stored at pointer p
<code>__m128i _mm_mul_ps(__m128 a, __m128 b)</code>	returns vector (a0*b0, a1*b1, a2*b2, a3*b3)
<code>void _mm_storeu_si128(__m128i *p, __m128i a)</code>	stores 128-bit vector a at pointer p

1. Implement the following function, which returns the sum of two arrays:

```
static int product_naive(int n, int *a) {  
    int product = 1;  
    for (int i = 0; i < n; i++) {  
        product *= a[i];  
    }  
    return product;  
}  
  
static int product_vectorized(int n, int *a) {  
    int result[4];  
    __m128i prod_v = __mm_load1_si128();  
  
    for (int i = 0; i < n/4 * 4; i += 4) { // Vectorised loop  
        prod_v = __mm_mul_ps(prod_v, __mm_loadu_si128((__m128i *) (a + i)));  
    }  
    _mm_storeu_si128((__m128i *) result, prod_v);  
  
    for (int i = n/4 * 4; i < n; i++) { // Handle tail case  
        result[0] *= a[i];  
    }  
    return result[0] * result[1] * result[2] * result[3];  
}
```