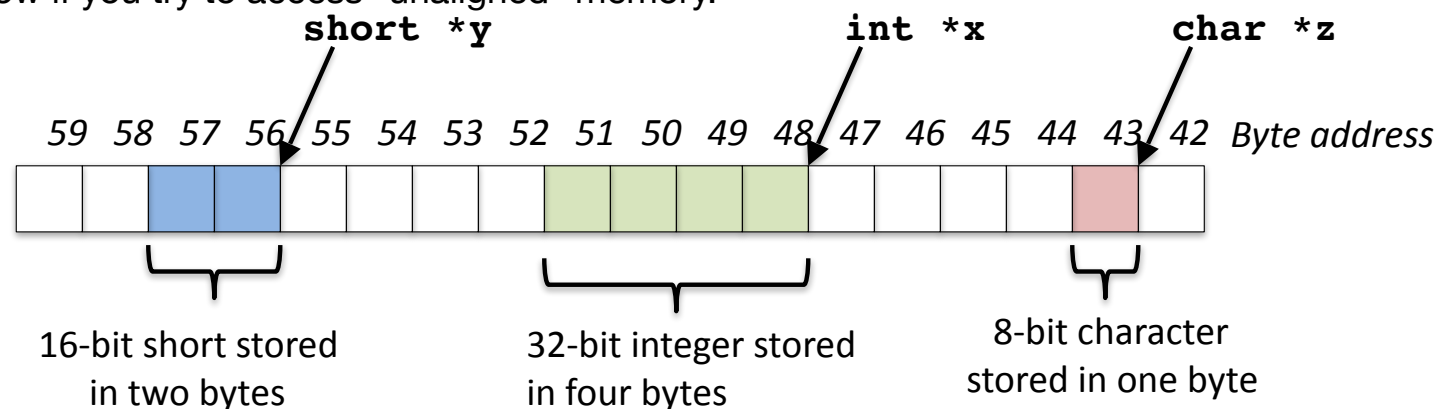


Arrays and Memory Management

Pointing to Different Size Objects

- Modern machines are “byte-addressable”
 - Hardware’s memory composed of 8-bit storage cells, each has a unique address
- A C pointer is just abstracted memory address
- Type declaration tells compiler how many bytes to fetch on each access through pointer
 - E.g., 32-bit integer stored in 4 consecutive 8-bit bytes
- But we actually want “word alignment”
 - Some processors will not allow you to address 32b values without being on 4 byte boundaries
 - Others will just be very slow if you try to access “unaligned” memory.



sizeof() operator

- **sizeof(type)** returns number of bytes in object
- But number of bits in a byte is not standardized
 - In olden times, when dragons roamed the earth, bytes could be 5, 6, 7, 9 bits long
- By Standard C99 definition, **sizeof(char) == 1**
- C does not play well with Unicode (unlike Python), so no `char c = '💩'`
- Can take **sizeof(arg)**, or **sizeof(structtype)**
- We'll see more of sizeof when we look at dynamic memory management
 - Remember, C **does not specify** **sizeof(int)**, so you need to ask it whenever you need to find this out!

Pointer Arithmetic

pointer + number *pointer – number*

e.g., *pointer + 1* adds 1 something to a pointer

```
char *p;  
char a;  
char b;  
  
p = &a;  
p += 1;
```

In each, p now points to b
(Assuming compiler doesn't
reorder variables in memory.

Never code like this!!!!

```
int *p;  
int a;  
int b;  
  
p = &a;  
p += 1;
```

Adds `1*sizeof(char)`
to the memory address

Adds `1*sizeof(int)`
to the memory address

Pointer arithmetic should be used cautiously

Valid Pointer Arithmetic

- Add an integer to a pointer.
- Subtracting an integer from a pointer
- Subtract 2 pointers (in the same array & of the same type)
- Compare pointers (<, <=, ==, !=, >, >=)
- Compare pointer to NULL (indicates that the pointer points to nothing)
- Everything else illegal since makes no sense:
 - adding two pointers
 - multiplying pointers
 - subtract pointer from integer

C Arrays

- Declaration:

```
int ar[2];
```

declares a 2-element integer array: just a block of memory which is uninitialized

```
int ar[] = {795, 635};
```

declares and initializes a 2-element integer array

Array Name / Pointer Duality

- *Key Concept:* Array variable is simply a “pointer” to the first (0th) element
- So, array variables almost identical to pointers
 - `char *string` and `char string[]` are nearly identical declarations
 - Differ in subtle ways: incrementing, declaration of filled arrays
- Consequences:
 - `ar` is an array variable, but works like a pointer
 - `ar[0]` is the same as `*ar`
 - `ar[2]` is the same as `*(ar+2)`
 - Can use pointer arithmetic to access arrays

C Arrays are Very Primitive

- An array in C does not know its own length, ***and its bounds are not checked!***
 - Consequence: We can accidentally ***access off the end of an array***
 - Consequence: We must pass the array ***and its size*** to any procedure that is going to manipulate it
- Segmentation faults and bus errors:
 - These are VERY difficult to find;
be careful! (You'll learn how to debug these in lab)
 - But also “fun” to exploit:
 - “Stack overflow exploit”, maliciously write off the end of an array on the stack
 - “Heap overflow exploit”, maliciously write off the end of an array on the heap

C Strings

- String in C is just an array of characters

```
char string[] = "abc";
```

- How do you tell how long a string is?
- Last character is followed by a 0 byte (aka “null terminator”):
written as 0 (the number) or '\0'
as a character

```
int strlen(char s[])  
{  
    int n = 0;  
    while (s[n] != 0) {  
        n++;  
    }  
    return n;  
}
```

Use Defined Constants

- Array size n ; want to access from 0 to $n-1$, so you should use counter AND utilize a variable for declaration & incrementation
 - Bad pattern

```
int i, ar[10];
for(i = 0; i < 10; i++){ ... }
```
 - Better pattern

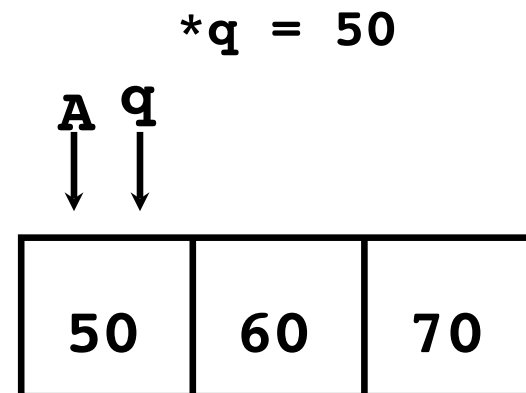
```
const int ARRAY_SIZE = 10;
int i, a[ARRAY_SIZE];
for(i = 0; i < ARRAY_SIZE; i++){ ... }
```
- **SINGLE SOURCE OF TRUTH**
 - You're utilizing indirection and avoiding maintaining two copies of the number 10
 - DRY: "Don't Repeat Yourself"
 - And don't forget the $<$ rather than $<=$:
When Nick took 60c, he lost a day to a "segfault in a malloc called by printf on large inputs":
Had a $<=$ rather than a $<$ in a single array initialization!

Changing a Pointer Argument?

- What if want function to change a pointer?
- What gets printed?

```
void inc_ptr(int *p)
{   p = p + 1;   }

int A[3] = {50, 60, 70};
int* q = A;
inc_ptr( q);
printf(" *q = %d\n", *q);
```

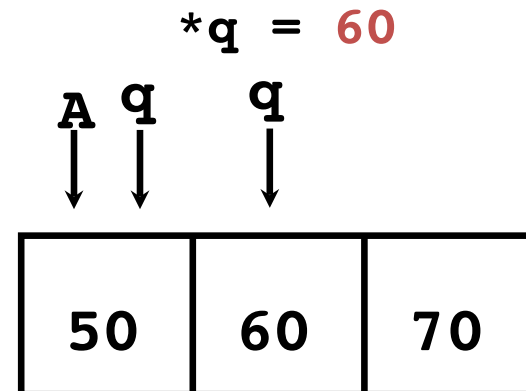


Pointer to a Pointer

- Solution! Pass a pointer to a pointer, declared as ****h**
- Now what gets printed?

```
void inc_ptr(int **h)
{    *h = *h + 1;    }

int A[3] = {50, 60, 70};
int* q = A;
inc_ptr(&q);
printf("*q = %d\n", *q);
```



When Arrays Go Bad: Heartbleed

- In TLS encryption, messages have a length...
 - And get copied into memory before being processed
- One message was “Echo Me back the following data, its this long...”
 - But the (different) echo length wasn’t checked to make sure it wasn’t too big...

```
M 5 HB L=5000 107:Ou17;GET / HTTP/1.1\r\n
Host: www.mydomain.com\r\nCookie: login=1
17kf9012oeu\r\nUser-Agent: Mozilla...
```

- So you send a small request that says “read back a lot of data”
 - And thus get web requests with auth cookies and other bits of data from random bits of memory...

Arrays and Pointers

- Array \approx pointer to the initial element
 - $a[i] \equiv *(a+i)$
- An array is passed to a function as a pointer
 - The array size is lost!
- Usually bad style to interchange arrays and pointers
 - Avoid pointer arithmetic!
 - Especially avoid things like `ar++`;

Passing arrays:

Really `int *array` Must explicitly pass the size

```
int
foo(int array[],
    unsigned int size)
{
    ... array[size - 1] ...
}

int
main(void)
{
    int a[10], b[5];
    ... foo(a, 10) ... foo(b, 5) ...
}
```

Arrays and Pointers

```
int
foo(int array[],
    unsigned int size)
{
    ...
    printf("%d\n", sizeof(array));
}

int
main(void)
{
    int a[10], b[5];
    ... foo(a, 10) ... foo(b, 5) ...
    printf("%d\n", sizeof(a));
}
```

What does this print? **4**

... because **array** is really a pointer (and a pointer is architecture dependent, but likely to be 4 or 8 on modern 32-64 bit machines!)

What does this print? **40**

Arrays and Pointers

```
int i;  
int array[10];  
  
for (i = 0; i < 10; i++)  
{  
    array[i] = ...;  
}
```

```
int *p;  
int array[10];  
  
for (p = array; p < &array[10]; p++)  
{  
    *p = ...;  
}
```

These code sequences have the same effect!

But the former is ***much more readable***:

Especially don't want to see code like **ar++**

Clickers/Peer Instruction Time

```
int x[] = { 2, 4, 6, 8, 10 };  
int *p = x;  
int **pp = &p;  
(*pp)++;  
(*(*pp))++;  
printf("%d\n", *p);
```

Result is:

A: 2

B: 3

C: 4

D: 5

E: None of the above

Clickers/Peer Instruction Time

```
int x[] = { 2, 4, 6, 8, 10 };  
int *p = x;  
int **pp = &p;  
(*pp)++;  
(*(*pp))++;  
printf("%d\n", *p);
```

P points to the start of X (2)

PP points to P

Increments P point to 2nd element (4)

Increments 2nd element by 1 (5)

Result is:

A: 2

B: 3

C: 4

D: 5

E: None of the above

Administrivia

- hw0 - edX is due on 2/13
- hw0 mini bio is due in second section
- Project 1 and the first lab will be out RSN
 - "Real Soon Now"
- MT1 will be Friday, 2/24 from 7 PM - 9 PM
- MT2 will be Thursday, 4/06 from 7 PM - 9 PM
 - Email the head Rebecca and Stephen for conflicts

Concise strlen()

```
int strlen(char *s)
{
    char *p = s;
    while (*p++)
        ; /* Null body of while */
    return (p - s - 1);
}
```

What happens if there is no zero character at end of string?

Arguments in `main()`

- To get arguments to the main function, use:
 - `int main(int argc, char *argv[])`
- What does this mean?
 - `argc` contains the number of strings on the command line (the executable counts as one, plus one for each argument). Here `argc` is 2:
 - `unix% sort myFile`
 - `argv` is a pointer to an array containing the arguments as strings
 - Since it is an array of pointers to character arrays
 - Sometimes written as `char **argv`

Example

- `foo hello 87 "bar baz"`
- `argc = 4 /* number arguments */`
- `argv[0] = "foo",`
`argv[1] = "hello",`
`argv[2] = "87",`
`argv[3] = "bar baz",`
- Array of pointers to strings

C Memory Management

- How does the C compiler determine where to put all the variables in machine's memory?
- How to create dynamically sized objects?
- To simplify discussion, we assume *one program runs at a time*, with access to all of memory.
- Later, we'll discuss ***virtual memory***, which lets multiple programs all run at same time, each thinking they own all of memory.

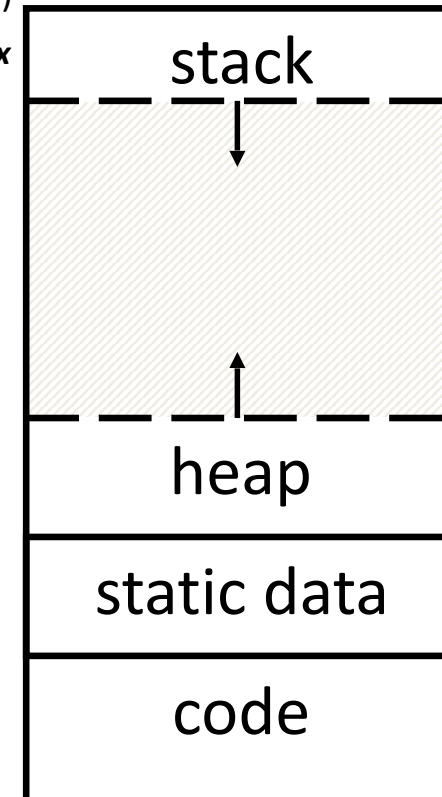
C Memory Management

- Program's address space contains 4 regions:
 - **stack**: local variables inside functions, grows downward
 - **heap**: space requested for dynamic data via `malloc()` resizes dynamically, grows upward
 - **static data**: variables declared outside functions, does not grow or shrink. Loaded when program starts, can be modified.
 - **code**: loaded when program starts, does not change

Memory Address

(32 bits assumed here)

$\sim FFFF\ FFFF_{hex}$



$\sim 0000\ 0000_{hex}$

Where are Variables Allocated?

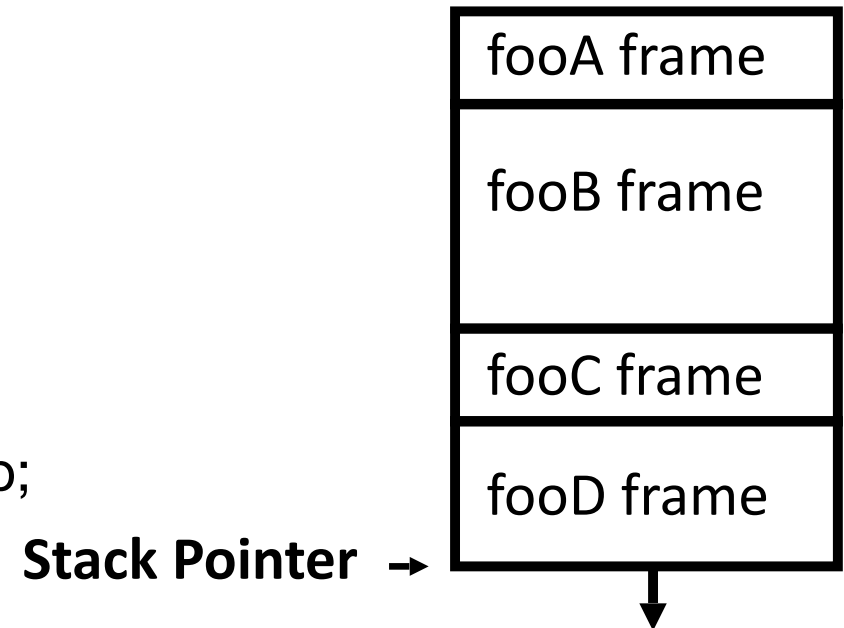
- If declared outside a function, allocated in “static” storage
- If declared inside function, allocated on the “stack” and freed when function returns
- `main()` is treated like a function
- For both of these types of memory, the management is automatic:
 - You don't need to worry about deallocating when you are no longer using them

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

The Stack

- Every time a function is called, a new frame is allocated on the stack
- Stack frame includes:
 - Return address (who called me?)
 - Arguments
 - Space for local variables
- Stack frames uses contiguous blocks of memory; stack pointer indicates start of stack frame
- When function ends, stack pointer moves up; frees memory for future stack frames
- We'll cover details later for MIPS processor

```
fooA() { fooB(); }  
fooB() { fooC(); }  
fooC() { fooD(); }
```



Stack Animation

- Last In, First Out (LIFO) data structure

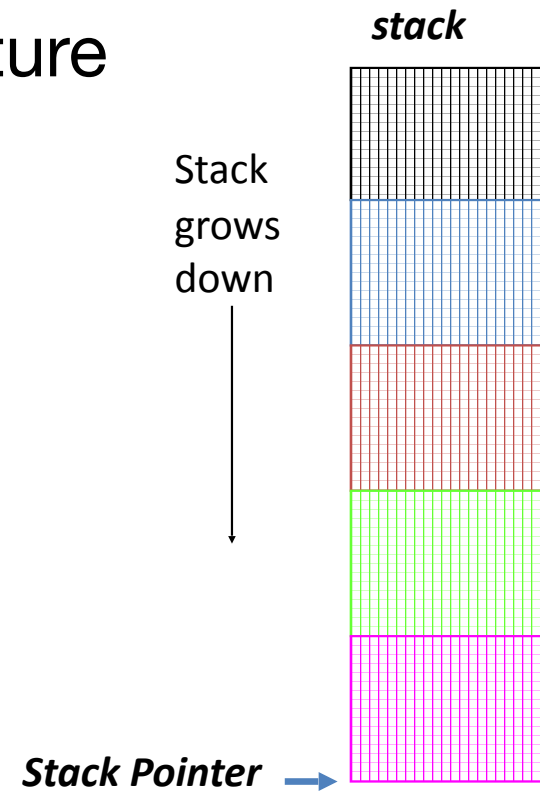
```
main ()
{ a(0);
}

void a (int m)
{ b(1);
}

void b (int n)
{ c(2);
}

void c (int o)
{ d(3);
}

void d (int p)
{
}
```



Managing the Heap

C supports functions for heap management:

- **malloc()** allocate a block of uninitialized memory
- **calloc()** allocate a block of zeroed memory
- **free()** free previously allocated block of memory
- **realloc()** change size of previously allocated block
 - careful – it might move!
 - And it ***will not update other pointers pointing to the same block of memory***

Malloc()

- **`void *malloc(size_t n):`**
 - Allocate a block of uninitialized memory
 - NOTE: Subsequent calls probably will not yield adjacent blocks
 - **`n`** is an integer, indicating size of requested memory block in bytes
 - **`size_t`** is an unsigned integer type big enough to “count” memory bytes
 - Returns **`void*`** pointer to block; **`NULL`** return indicates no more memory (check for it!)
 - Additional control information (including size) stored in the heap for each allocated block.
- Examples:
 - *“Cast” operation, changes type of a variable.*
*Here changes (**`void *`**) to (**`int *`**)*

```
int *ip;  
ip = (int *) malloc(sizeof(int));
```
 - ```
typedef struct { ... } TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
```
- **`sizeof`** returns size of given type in bytes, ***necessary if you want portable code!***

# And then free()

- **void free(void \*p) :**
  - **p** is a pointer containing the address originally returned by **malloc()**
- Examples:
  - ```
int *ip;  
ip = (int *) malloc(sizeof(int));  
... ..  
free((void*) ip); /* Can you free(ip) after ip++ ? */
```
 - ```
typedef struct {... } TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
... ..
free((void *) tp);
```
- When you free memory, you must be sure that you pass the original address returned from **malloc()** to **free()**; Otherwise, crash (or worse)!

# Using Dynamic Memory

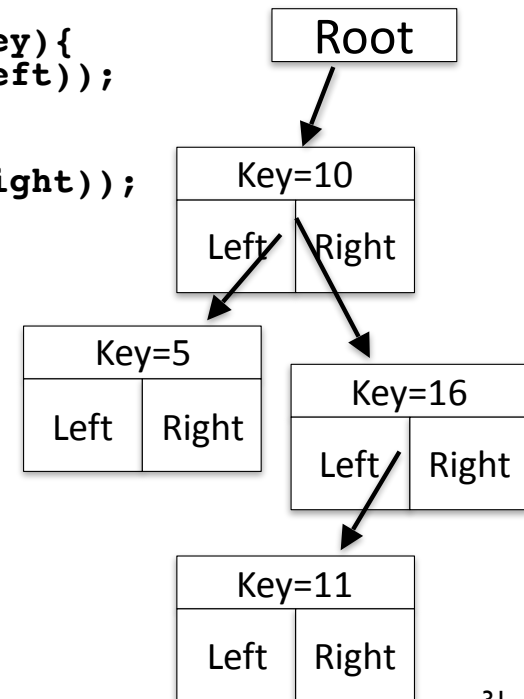
```
typedef struct node {
 int key;
 struct node *left; struct node
 *right;
} Node;

Node *root = NULL;

Node *create_node(int key, Node
 *left,
 Node *right){
 Node *np;
 if(!(np =
 (Node*) malloc(sizeof(Node))){
 printf("Memory exhausted!\n");
 exit(1);}
 else{
 np->key = key;
 np->left = left;
 np->right = right;
 return np;
 }
}
```

```
void insert(int key, Node **tree){
 if ((*tree) == NULL){
 (*tree) = create_node(key, NULL,
 NULL);
 }
 else if (key <= (*tree)->key){
 insert(key, &((*tree)->left));
 }
 else{
 insert(key, &((*tree)->right));
 }
}

int main(){
 insert(10, &root);
 insert(16, &root);
 insert(5, &root);
 insert(11, &root);
 return 0;
}
```



# Observations

- Code, Static storage are easy: they never grow or shrink
- Stack space is relatively easy: stack frames are created and destroyed in last-in, first-out (LIFO) order
- Managing the heap is tricky: memory can be allocated / deallocated at any time
  - If you forget to deallocate memory: “Memory Leak”
    - Your program ***will eventually run out of memory***
  - If you call free twice on the same memory: “Double Free”
    - Possible ***crash or exploitable vulnerability***
  - If you use data after calling free: “Use after free”
    - Possible ***crash or exploitable vulnerability***



# And In Conclusion, ...

- C has three main memory segments in which to allocate data:
  - Static Data: Variables outside functions
  - Stack: Variables local to function
  - Heap: Objects explicitly malloc-ed/free-d.
- Heap data is biggest source of bugs in C code