

Lecture 6



Computer Science 61C Spring 2017

January 30th, 2017

Friedland and Weaver

Memory Management and more

Administrivia

- My office hours: Monday 1pm-2pm, 424 SDH.
- Raspberry PI servers online today!

Agenda

- Memory Management
- and more

C Memory Management

- How does the C compiler determine where to put all the variables in machine's memory?
- How to create dynamically sized objects?
- To simplify discussion, we assume *one program runs at a time*, with access to all of memory.
- Later, we'll discuss ***virtual memory***, which lets multiple programs all run at same time, each thinking they own all of memory.


C Memory Management

- Program's address space contains 4 regions:
 - **stack**: local variables inside functions, grows downward
 - **heap**: space requested for dynamic data via `malloc()` resizes dynamically, grows upward
 - **static data**: variables declared outside functions, does not grow or shrink. Loaded when program starts, can be modified.
 - **code**: loaded when program starts, does not change

Memory Address

(32 bits assumed here)

$\sim FFFF\ FFFF_{hex}$



$\sim 0000\ 0000_{hex}$

Where are Variables Allocated?


- If declared outside a function, allocated in “static” storage
- If declared inside function, allocated on the “stack” and freed when function returns
- `main()` is treated like a function
- For both of these types of memory, the management is automatic:
 - You don't need to worry about deallocating when you are no longer using them

```
int myGlobal;  
main() {  
    int myTemp;  
}
```

The Stack

- Every time a function is called, a new frame is allocated on the stack
- Stack frame includes:
 - Return address (who called me?)
 - Arguments
 - Space for local variables
- Stack frames use contiguous blocks of memory; stack pointer indicates start of stack frame
- When function ends, stack pointer moves up; frees memory for future stack frames
- We'll cover details later for MIPS processor

```
fooA() { fooB(); }  
fooB() { fooC(); }  
fooC() { fooD(); }
```



Stack Animation

- Last In, First Out (LIFO) data structure


```
main ()
{ a(0);
}

void a (int m)
{ b(1);
}

void b (int n)
{ c(2);
}

void c (int o)
{ d(3);
}

void d (int p)
{
}
```



Managing the Heap

C supports functions for heap management:

- **malloc()** allocate a block of uninitialized memory
- **calloc()** allocate a block of zeroed memory
- **free()** free previously allocated block of memory
- **realloc()** change size of previously allocated block
 - careful – it might move!
 - And it ***will not update other pointers pointing to the same block of memory***

Malloc()

- **`void *malloc(size_t n):`**
 - Allocate a block of uninitialized memory
 - NOTE: Subsequent calls probably will not yield adjacent blocks
 - **`n`** is an integer, indicating size of requested memory block in bytes
 - **`size_t`** is an unsigned integer type big enough to “count” memory bytes
 - Returns **`void*`** pointer to block; **`NULL`** return indicates no more memory (check for it!)
 - Additional control information (including size) stored in the heap for each allocated block.
- Examples:
 - *“Cast” operation, changes type of a variable.
Here changes **`(void *)`** to **`(int *)`***

```
int *ip;  
ip = (int *) malloc(sizeof(int));
```
 - ```
typedef struct { ... } TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
```
- **`sizeof`** returns size of given type in bytes, ***necessary if you want portable code!***

# And then free()

- **void free(void \*p) :**
  - **p** is a pointer containing the address originally returned by **malloc()**
- Examples:
  - ```
int *ip;  
ip = (int *) malloc(sizeof(int));  
... ..  
free((void*) ip); /* Can you free(ip) after ip++ ? */
```
 - ```
typedef struct {... } TreeNode;
TreeNode *tp = (TreeNode *) malloc(sizeof(TreeNode));
... ..
free((void *) tp);
```
- When you free memory, you must be sure that you pass the original address returned from **malloc()** to **free()**; Otherwise, crash (or worse)!

# Using Dynamic Memory


```
typedef struct node {
 int key;
 struct node *left; struct node
 *right;
} Node;

Node *root = NULL;

Node *create_node(int key, Node
 *left,
 Node *right){
 Node *np;
 if(!(np =
 (Node*) malloc(sizeof(Node)))){
 printf("Memory exhausted!\n");
 exit(1);}
 else{
 np->key = key;
 np->left = left;
 np->right = right;
 return np;
 }
}
```

```
void insert(int key, Node **tree){
 if ((*tree) == NULL){
 (*tree) = create_node(key, NULL,
 NULL);
 }
 else if (key <= (*tree)->key){
 insert(key, &((*tree)->left));
 }
 else{
 insert(key, &((*tree)->right));
 }
}

int main(){
 insert(10, &root);
 insert(16, &root);
 insert(5, &root);
 insert(11, &root);
 return 0;
}
```



# Observations

- Code, Static storage are easy: they never grow or shrink
- Stack space is relatively easy: stack frames are created and destroyed in last-in, first-out (LIFO) order
- Managing the heap is tricky: memory can be allocated / deallocated at any time
  - If you forget to deallocate memory: “Memory Leak”
    - Your program ***will eventually run out of memory***
  - If you call free twice on the same memory: “Double Free”
    - Possible ***crash or exploitable vulnerability***
  - If you use data after calling free: “Use after free”
    - Possible ***crash or exploitable vulnerability***

# And In Conclusion, ...

- C has three main memory segments in which to allocate data:
  - Static Data: Variables outside functions
  - Stack: Variables local to function
  - Heap: Objects explicitly malloc-ed/free-d.
- Heap data is biggest source of bugs in C code

## Clickers/Peer Instruction!

```
int x = 2;
int result;

int foo(int n)
{ int y;
 if (n <= 0) { printf("End case!\n"); return 0; }
 else
 { y = n + foo(n-x);
 return y;
 }
}
result = foo(10);
```

Right after the **printf** executes but before the **return 0**, how many copies of **x** and **y** are there allocated in memory?

- A: #x = 1, #y = 1
- B: #x = 1, #y = 5
- C: #x = 5, #y = 1
- D: #x = 1, #y = 6
- E: #x = 6, #y = 6